

SOHAM CHAUDHARI

+91 7499787431 — Mumbai, Maharashtra, India — sohamrc08@gmail.com — [LinkedIn](#) — [GitHub](#)

SUMMARY

AI/ML Engineer with production experience architecting Agentic AI systems, Retrieval-Augmented Generation (RAG) pipelines, and multimodal machine learning solutions across fintech and legal-tech. Delivered measurable impact through LLM optimization, hybrid vector search, and high-concurrency API design, including a 60% gain in retrieval accuracy and a 60–70% reduction in token usage. Specializing in Generative AI, Natural Language Processing, and Computer Vision.

EXPERIENCE

AI Intern

Feb 2026 – Present

FLIP (paybyflip)

Remote

- Engineered an 11-node LangGraph orchestration pipeline powering an AI-driven credit card recommendation platform end-to-end
- Designed a 4-step hybrid retrieval engine combining SQL filtering, 1536-dimensional vector embeddings, semantic search, and Tavily web fallback across 9.7K+ offer entries
- Implemented PostgreSQL-backed semantic memory with async storage, session isolation, and feedback APIs; optimized memory injection to the top 1–2 items, reducing token usage by 60–70%
- Deployed a production-ready backend on AWS EC2 via Docker Compose with connection pooling (10–50 DB connections), request timeouts, health/readiness endpoints, and API rate limiting at 100 req/min

AI Developer Intern

Sep 2025 – Feb 2026

UnLawC

Remote

- Delivered end-to-end Legal Operations AI agents powered by LLMs, cutting manual document review time by 50%
- Optimized LLM performance through prompt engineering, improving contextual accuracy by 31%
- Architected RAG pipelines using Mistral AI, LangChain, and Pinecone, boosting retrieval accuracy by 60% and eliminating hallucinations

EDUCATION

Bachelor of Engineering, Computer Engineering

Nov 2022 – Expected 2026

Vivekanand Education Society's Institute of Technology — CGPA: 8.27/10.0

PROJECTS

MOSAIC – Multimodal Video Analysis Platform

[GitHub](#)

- Built a multimodal search engine using CLIP and FAISS vector search, achieving 99.9% recall for natural language queries
- Designed agentic workflows with LangGraph, enabling autonomous video segmentation with 90% precision
- Enhanced video processing throughput using asyncio and FFmpeg, supporting 5+ concurrent streams with sub-500ms latency

VISION AI – Deep Learning Super-Resolution System

[GitHub](#)

- Developed an ESRGAN-based super-resolution model, achieving 4x resolution enhancement with 92% quality improvement
- Trained a Generative Adversarial Network (GAN) using PyTorch on the DIV2K dataset, with hyperparameter tuning via Optuna
- Launched a production backend on Render and frontend on Vercel for real-time image processing

TECHNICAL SKILLS

Programming Python, JavaScript, Java, SQL

AI & LLMs LangChain, LangGraph, RAG, Generative AI, Prompt Engineering, OpenAI, Claude, Mistral

ML & DL PyTorch, TensorFlow, Keras, scikit-learn, Computer Vision, NLP, CLIP, Transformers

Databases FAISS, ChromaDB, Pinecone, Weaviate, MongoDB, PostgreSQL

MLOps Docker, Kubernetes, MLflow, CI/CD, Git, GitHub Actions, AWS

Web & APIs FastAPI, Flask, Node.js, Express.js, React.js, Next.js